

A Sparse Probing Results




Figure 4: Sparse probing results. In these experiments, we train the same probes as in the main paper, but with a varying ℓ_1 penalty applied to the probing objective to encourage the discovery of sparse solutions. Here we report the level of sparsity, probe accuracy, and distribution of disagreement types as we vary the strength of the regularizer within $\{0, 0.01, 0.03, 0.1\}$. Except at extremely high sparsity values, both accuracies and error distributions remain similar to those reported in the main set of experiments.