

Model Performance of 1st-Person vs. 3rd-Person Prompts on 3 Different Expertise Levels




Figure 14: First-pov consistently produces more sycophancy than Third-pov across all 3 expertise levels. This is a full figure for Figure 8.

Cosine Similarity between Category Centroids across 1st- and 3rd-person Prompts




Figure 15: Cosine similarity heatmap between category centroids obtained from PCA-projected hidden states at layer 32 of *Qwen2.5 7B-Instruct*. Each centroid represents a specific combination of pronoun perspective (1st vs. 3rd person) and expertise level (Beginner, Intermediate, Advanced).

first- vs. third-person (perspective) prompts are more distinct. This confirms that perspective cues have a stronger effect on model representations and sycophancy than authority cues, generalizing across model families.